

16-720 Computer Vision: Homework 5

Optical Character Recognition using Neural Networks

Instructors: Deva Ramanan

TAs: Achal Dave*, Shashank Jujjavarapu, Siddarth Malreddy, Brian Pugh

Originally developed by: Syed Bokhari, Rahul Nallamothu

- Please pack your system and write-up into a single file `<andrewid>.zip`, in accordance with the complete submission checklist at the end of this document.
- All tasks marked with a **Q** require a submission.
- Please stick to the provided function signatures, variable names, and file names.
- **Start early!** This homework cannot be completed within two hours!
- **Verify your implementation as you proceed** otherwise you will risk having a huge mess of malfunctioning code that can go wrong anywhere.
- If you have any questions, please post them to Piazza, or come to office hours.
- Read through the whole assignment once before starting.



Figure 1: Samples from NIST Special 19 dataset [1]

1 Overview

Deep learning has quickly become one of the most applied machine learning techniques in computer vision. Convolutional neural networks have been applied to many different computer vision problems such as image classification, recognition, and segmentation with great success. In this assignment, you will first implement a fully connected feed forward neural network for hand written character classification. Then in the second part, you will implement a system to locate characters in an image, which you can then classify with your deep network. The end result will be a system that, given an image of hand written text, will output the text contained in the image.

1.1 Basic Use

Here we will give a brief overview of the math for a single hidden layer feed forward network. For a more detailed look at the math and derivation, please see the class slides.




Figure 2: Example of a single hidden layer network

A fully-connected network, \mathbf{f} , for classification, applies a series of linear and non-linear functions to an input data vector \mathbf{x} of size $N \times 1$ to produce an output vector $\mathbf{f}(\mathbf{x})$ of size $C \times 1$, where each element i of the output vector represents the probability of \mathbf{x} belonging to the class i . Since the data samples are of dimensionality N , this means the input layer

has N input units. To compute the value of the output units, we must first compute the values of all the hidden layers. The first hidden layer *pre-activation* $\mathbf{a}^{(1)}(\mathbf{x})$ is given by

$$\mathbf{a}^{(1)}(\mathbf{x}) = \mathbf{W}^{(1)}\mathbf{x} + \mathbf{b}^{(1)}$$

Then the *post-activation* values of the first hidden layer $\mathbf{h}^{(1)}(\mathbf{x})$ are computed by applying a non-linear activation function \mathbf{g} to the *pre-activation* values

$$\mathbf{h}^{(1)}(\mathbf{x}) = \mathbf{g}(\mathbf{a}^{(1)}(\mathbf{x})) = \mathbf{g}(\mathbf{W}^{(1)}\mathbf{x} + \mathbf{b}^{(1)})$$

The pre- and post- activations of Subsequent hidden layers ($1 < t \leq T$) are given by:

$$\begin{aligned}\mathbf{a}^{(t)}(\mathbf{x}) &= \mathbf{W}^{(t)}\mathbf{h}^{(t-1)} + \mathbf{b}^{(t)} \\ \mathbf{h}^{(t)}(\mathbf{x}) &= \mathbf{g}(\mathbf{a}^{(t)}(\mathbf{x}))\end{aligned}$$

The output layer *pre-activations* $\mathbf{a}^{(T)}(\mathbf{x})$ are computed in a similar way

$$\mathbf{a}^{(T)}(\mathbf{x}) = \mathbf{W}^{(T)}\mathbf{h}^{(T-1)}(\mathbf{x}) + \mathbf{b}^{(T)}$$

and finally the *post-activation* values of the output layer are computed with

$$\mathbf{f}(\mathbf{x}) = \mathbf{o}(\mathbf{a}^{(T)}(\mathbf{x})) = \mathbf{o}(\mathbf{W}^{(T)}\mathbf{h}^{(T-1)}(\mathbf{x}) + \mathbf{b}^{(T)})$$

where \mathbf{o} is the output activation function.

For this assignment, we will be using the sigmoid activation function for the hidden layer, so:

$$\mathbf{g}(y) = \frac{1}{1 + \exp(-y)}$$

where when \mathbf{g} is applied to a vector, it is applied element wise across the vector.

Since we are using this network for classification, a common output activation function to use is the softmax function. This will allow us to map the real valued activations $\mathbf{a}^{(T)}(\mathbf{x})$ into a probability distribution (vector of positive numbers that sum to 1). Letting \mathbf{x}_i denote the i^{th} element of the vector \mathbf{x} , the softmax function is defined as:

$$\mathbf{o}_i(\mathbf{y}) = \frac{\exp(\mathbf{y}_i)}{\sum_j \exp(\mathbf{y}_j)}$$

Q1.1.1 “Theory” [5 points] In training deep networks, the ReLU activation function is generally preferred to the sigmoid activation function. Why might this be the case?

Q1.1.2 Theory [5 points] All types of deep networks use non-linear activation functions for their hidden layers. Suppose we have a neural network with input dimension N and output dimension C and T hidden layers. Prove that if we have a linear activation function \mathbf{g} , then the number of hidden layers has no effect on the representation capability of the network (i.e., that the set of functions that can be represented by a T layer network is exactly the same as the set that can be represented by a $T' \neq T$ layer network).

2 Implement a Fully Connected Network

In this section, you will implement all of the functions needed to initialize, train, evaluate, and use the network. Throughout this assignment, the network will be represented by its parameters \mathbf{W} and \mathbf{b} , the weights and biases of the network.

2.1 Network Initialization

In order to use a deep network, we must first initialize the weights and biases in the network. This is typically done with a random initialization, or initializing the weights from some other training procedure (which we will see in Section 3). If you are interested in issues relating to initialization, [2] is a good paper to look at.

Q2.1.1 Theory [5 points] Why is it not a good idea to initialize a network with all zeros? How about all ones, or some other constant value? (Hint: Consider what the gradients from backpropagation will look like.)

Q2.1.2 Code [5 points] Implement `[W, b] = InitializeNetwork(layers)`. This function should take as input the sizes of the layers for your neural network. The `layers` parameters will be a vector of at least 3 integers. The first element denotes the size of the data layer N , the last element denotes the size of the output layer/number of classes C , and the intermediate elements denote the size of the hidden layers. You should return \mathbf{W} and \mathbf{b} which are both cell arrays containing `length(layers)-1` elements. \mathbf{W} should contain at least two weight matrices of appropriate size, and \mathbf{b} should contain at least two bias vectors.

Q2.1.3 Writeup [5 points] Describe the initialization you implemented in Q2.1.2 and any reasoning behind why you chose that strategy.

2.2 Forward Propagation

Section 1 has the math for forward propagation. The loss function generally used for classification is the cross-entropy loss.

$$L_{\mathbf{f}}(\mathbf{D}) = - \sum_{(\mathbf{x}, \mathbf{y}) \in \mathbf{D}} \log(\mathbf{y} \cdot \mathbf{f}(\mathbf{x}))$$

Here \mathbf{D} is the full training dataset of data samples \mathbf{x} ($N \times 1$ vectors) and labels \mathbf{y} ($C \times 1$ one-hot vectors).

Q2.2.1 Code [10 points] Implement `[out, act_a, act_h] = Forward(W, b, X)` runs forward propagation on an input data sample \mathbf{X} , using the network defined by the parameters \mathbf{W} and \mathbf{b} . Both \mathbf{W} and \mathbf{b} are cell arrays containing the network parameters as initialized by `InitializeNetwork(...)`. \mathbf{X} is a vector of dimensions $N \times 1$. This function should return the final softmax output of the network in the variable `out`, which is of size $C \times 1$. The function must also return the cell arrays `act_a` and `act_h`, which contain the pre- and post-activations of the network on this input sample.

Q2.2.2 Code [5 points] Implement `[outputs] = Classify(W, b, data)`. This function should accept the network parameters `W` and `b`, as well as a $D \times N$ matrix of data samples. `outputs` should contain the softmax output of the network for each of the data samples.

Q2.2.3 Code [10 points] Implement `[accuracy, loss] = ComputeAccuracyAndLoss(W, b, data, labels)`. This function should compute the accuracy on the network with respect to the provided `data` matrix of size $D \times N$ and `labels` matrix of size $D \times C$. You should also compute and return the average cross-entropy loss on the dataset.

2.3 Backwards Propagation

Gradient descent is an iterative optimisation algorithm, used to find the local optima. To find the local minima, we start at a point on the function and move in the direction of negative gradient (steepest descent) till some stopping criteria is met.

The update equation for a general weight $W_{ij}^{(t)}$ and bias $b_i^{(t)}$ is

$$W_{ij}^{(t)} = W_{ij}^{(t)} - \alpha * \frac{\partial L_{\mathbf{f}}}{\partial W_{ij}^{(t)}}(\mathbf{x}) \quad b_i^{(t)} = b_i^{(t)} - \alpha * \frac{\partial L_{\mathbf{f}}}{\partial b_i^{(t)}}(\mathbf{x})$$

α is the learning rate. Please refer to the class slides for more details on how to derive the gradients. Note that here we are using softmax loss (which is different from the derivations we saw in class).

Q2.3.1 Code [30 points] Implement `[grad_W, grad_b] = Backward(W, b, X, Y, act_h, act_a)` that runs the back propagation algorithm on a single input example `X` and the ground truth vector `Y`. The function takes as input the network parameters `W` and `b` and the network pre- and post-activations `act_a` and `act_h` from calling `Forward` on `X`. The function should return the computed gradient updates for the network parameters. `grad_W` and `grad_b` are both cell arrays identical in shape to the kinds produced by `InitializeNetwork(...)`, but containing the gradient updates to make to each of the network parameters.

Q2.3.2 Code [5 points] Implement `[W, b] = UpdateParameters(W, b, grad_W, grad_b, learning_rate)` that computes and returns the updated network parameters `W` and `b`. The function is given the old parameters, the parameters gradients `grad_W` and `grad_b` returned by `Backward(...)`, and the supplied `learning_rate`.

2.4 Training Loop

When training a neural network using Gradient Descent, there are two options one can take for the updates, batch or stochastic updates. In batch gradient descent, we compute the gradient for all the examples in the training set and then average the gradients before updating the weights. In stochastic, we compute the gradient with one sample, update the weights with that gradient and carry on with the next sample. Each update is called a *iteration*, one complete pass through the entire dataset is called an *epoch*. In case of stochastic gradient descent, number of iterations in an epoch equals the number of training samples.

Q2.4.1 Theory [5 points] Give pros and cons for both stochastic and batch gradient descent. In general, which one is faster to train in terms of number of epochs? Which one is faster in terms of number of iterations?

Q2.4.2 Code [10 points] Implement

`[W, b] = Train(W, b, train_data, train_label, learning_rate)` which trains the network for one epoch on `train_data` and ground-truth `train_label`. The function should return the updated network parameters.

2.5 Gradient Checker [25 points]

Often, when implementing new layers in popular packages like Caffe or Torch you will have to implement the forward and backward pass for that layer. While the forward is fairly easy, a lot of things can go wrong in the backward pass. One common way to check if the backward pass is correct is by writing a gradient checker. A good explanation of a gradient checker is in the link: <http://ufldl.stanford.edu/tutorial/supervised/DebuggingGradientChecking/>

Q2.3.3 Code [25 points] Implement a script `checkGradient.m` that checks gradients of the loss with respect to a few random weights in each layer. Submit the code.

3 Training Models

We now have all the code required to train and use a deep network. In this section, you will train and evaluate different models which will then be used in Section 4 for parsing text in images.

3.1 From Scratch

Often times when you encounter new problems in deep learning, you will need to train a new network from scratch. This means that you will randomly initialize the weights in some way, and run multiple iterations of back propagation. The `Train.m` function you implemented in runs one epoch of training, passing over all data samples once. To fully train a deep network you typically must do many epochs.

Once a deep network is trained, one way to gain some insight into what it is doing is to visualize the learned weights. With convolutional neural networks as seen in class, each of the learned filters can be visualized as an image. This can also be done with a fully connected network, as you have implemented here. Since our input images are 32×32 images, unrolled into one 1024 dimensional vector that gets multiplied by $\mathbf{W}^{(1)}$, each row of $\mathbf{W}^{(1)}$ can be seen as a weight image. Reshaping each row into a 32×32 image can give us an idea of what types of images each unit in the hidden layer has a high response to.

We have provided you three data `.mat` files to use for this section. The training data in `nist26_train.mat` contains 299 samples for each of the 26 upper-case letters of the alphabet. This is the set you should use for training your network. The cross-validation set in `nist26_valid.mat` contains 100 samples from each class, and should be used in the training




Figure 3: Samples weights from a network with 50 hidden units trained for 30 epochs. Your weights will likely look very different, as you are training with more hidden units.

loop to see how the network is performing on data that it is not training on. This will help to spot over fitting. Finally, the test data in `nist26_test.mat` contains another 100 samples per class, and should be used for the final evaluation on your best model to see how well it will generalize to new unseen data.

Q3.1.1 Code [5 points] Use the provided `train26.m` to train a network from scratch. Use a single hidden layer with 400 hidden units, and train for at least 30 epochs. **Modify** the script to plot generate two plots: one showing the accuracy on both the training and validation set over the epochs, and the other showing the cross-entropy loss. The x-axis should represent the epoch number, while the y-axis represents the accuracy or loss. With these settings, you should see an accuracy on the validation set of at least 75%.

Q3.1.2 Writeup [5 points] Use your modified training script to train two networks, one with learning rate 0.01, and another with learning rate 0.001. Include all 4 plots in your writeup. Comment on how the learning rates affect the training, and report the final accuracy of the best network on the test set.

Q3.1.3 Writeup [5 points] Using the best network from the previous question, report the accuracy and cross-entropy loss on the test set, and visualize the first layer weights that your network learned (using `reshape` and `montage`). Compare these to the network weights immediately after initialization. Include both visualizations in your writeup. Comment on the learned weights. Do you notice any patterns?

Q3.1.4 Writeup [5 points] Visualize the confusion matrix for your best model as a 26×26 image (upscale the image so we can actually see it). Comment on the top two pairs of classes that are most commonly confused.




Figure 4: Learned filters from AlexNet trained on the ImageNet dataset. AlexNet is a convolutional neural network whose architecture and weights are often used as the basis for new networks.

3.2 Fine Tuning

When training from scratch, a lot of epochs are often needed to learn anything meaningful. One way to avoid this is to instead initialize the weights more intelligently. Various strategies have been used for initializing neural networks, such as unsupervised pretraining with Auto-Encoders or Restricted Boltzmann Machines.

However thanks to the explosion in popularity of deep learning for computer vision, it is often possible to also initialize a network with weights from another deep network that was trained for a different purpose. This is because, whether we are doing image classification, segmentation, recognition etc..., most real images share common properties. Simply copying the weights from the other network to yours gives your network a head start, so your network does not need to learn these common weights from scratch all over again. This is also referred to as fine tuning.

We have trained a network for recognizing capital letters using 800 hidden units, and trained it for 60 epochs. The network parameters are stored in `nist26_model_60iters.mat`. Using these weights, you will initialize and fine-tune a network for recognizing both capital letters as well as numeric digits.

Q3.2.1 Code/Writeup [10 points] Make a copy of `train26.m` and name it `finetune36.m`. Modify this script to load the data from `nist36_*.mat`, and train a network to classify both written letters and numbers. Finetune (train) this network for 5 epochs with learning rate

0.01, and include plots of the accuracy and cross-entropy loss in your writeup.

Q3.2.2 Writeup [5 points] Once again, visualize the network's first layer weights before and after training. Comment on the differences you see. Also report the network's accuracy and loss on the test set.

Q3.2.3 Writeup [5 points] Visualize the confusion matrix for your best model as a 36×36 image (upscale the image so we can actually see it). Comment on the top two pairs of classes that are most commonly confused. How has introducing more classes affected the network?

4 Extract Text from Images




Figure 5: Sample image with handwritten characters annotated with boxes around each character.

Now that you have a network that can recognize handwritten letters with reasonable accuracy, you can now use it to parse text in an image. Given an image with some text on it, our goal is to have a function that returns the actual text in the image. However since your neural network expects a binary image with a single character, you will need to process the input image to extract each character. There are various approaches that can be done so feel free to use any strategy you like.

Here we outline one possible method:

1. Process the image (blur, threshold, etc...) to classify all pixels as being part of a character or background.
2. Find connected groups of character pixels (see `bwconncomp`). Place a bounding box around each connected component.

3. Group the letters based on which line of the text they are a part of, and sort each group so that the letters are in the order they appear on the page.
4. Take each bounding box one at a time and resize it to 32×32 , classify it with your network, and report the characters in order (inserting spaces when it makes sense).

Since the network you trained likely does not have perfect accuracy, you can expect there to be some errors in your final text parsing. Whichever method you choose to implement for the character detection, you should be able to place a box on most of these characters in the image. We have provided you with `01_list.jpg`, `02_letters.jpg`, `03_haiku.jpg` and `04_deep.jpg` to test your implementation on.

Q4.1 Theory [5 points] The method outlined above is pretty simplistic, and makes several assumptions. What are two big assumptions that the sample method makes. In your writeup, include two example images where you expect the character detection to fail (either miss valid letters, or respond to non-letters).

Q4.2 Code [25 points] Implement `[lines, bw] = findLetters(im)`. Given an RGB image, this function should return a cell array `lines` containing all of the located handwritten characters in the image, as well as a binary black-and-white version of the image `im`. The `letters` cell array should contain one entry for each line of text that appears in the image. Each entry should be a matrix of size $L_i \times 4$, where L_i represents the number of characters in the i^{th} line of text. Each row of the matrix should contain `[x1, y1, x2, y2]` the positions of the top-left and bottom-right corners of the box. The processed image `bw` should be only black and white (containing pixel values 0.0 and 1.0), with the characters in black and the rest of the image in white. Also include a test script `testFindLetters.m` that runs your function on all of the provided images (and any others you decided to include), and displays a figure showing the bounding boxes around each letter.

Q4.3 Writeup [5 points] Run `findLetters(...)` on all of the provided sample images in `images/`. Plot all of the located boxes on top of the image to show the accuracy of your `findLetters(...)` function. Include all the result images in your writeup.

Q4.4 Code [15 points] Implement `[text] = extractImageText(fname)`. This function should accept an image path `fname`, and return the text contained in the image. This function will load the image, find the character locations, classify each one with the network you trained in **Q3.2.1**, and return the text contained in the image. Also include a test script `testExtractImageText.m` that runs your function on all of the provided images (and any others you decided to include), and outputs the retrieved text for each image.

Q4.5 Writeup [5 points] Run your `extractImageText(...)` on all of the provided sample images in `images/`. Include the extracted text in your writeup.

5 Extra Credit

If you attempt any extra credit, put all code in the `ec` directory in your handin zip. Include a `README` describing how to run your extra credit. Additionally, include a paragraph in the writeup about what you attempted, as well as any relevant images or results.

5.1 Optimization Algorithms [10 points]

In Q2.4, we considered batch vs. stochastic gradient descent. In practice, there are a number of variants of gradient descent aimed at faster convergence or fewer hyperparameters. Implement one of the following variants: Adagrad [3], Adadelta [4], Adam [5], or RMSProp. For partial credit (5pts), you can instead implement stochastic gradient descent with momentum.

You may wish to look at the following link for some helpful details: <http://sebastianruder.com/optimizing-gradient-descent/index.html#rmsprop>.

5.2 Convolutional Neural Network [10 points]

In this assignment you implemented a simple fully connected feed forward neural network for character classification. Convolutional neural networks (CNNs) have proven to be far better suited to image related tasks. In this extra credit, you can use any popular framework (Caffe, Torch, TensorFlow, MatConvNet, Theano, MXNet, DL4J, MATLAB's deep learning toolbox, ConvNetJS, grenade, etc. - although we recommend one of the first few) to train a CNN to classify the handwritten characters. Describe the network architecture you used and include plots of the loss and accuracy.

5.3 Better OCR [20 points]

In Q4.1 you mentioned a few ways that the sample character detection algorithm would fail. Collect at least three such images (from the Internet or make them yourself) and write an more robust character detection algorithm. More points will be given for more robust approaches.

5.4 Dropout [10 points]

Dropout is a widely used regularization technique used while training deep networks. Re-implement the forward and backward propagation by incorporating dropout of 0.5. Compare the results in your report.

Note: You may not notice any increase in accuracy by using dropout for the 1 layer network since the capacity of the network is pretty limited.

5.5 Count the number of parameters [5 points]

Calculate the number of parameters in the VGG-16 network. The exact network we are interested in is listed as configuration D in on page 3 in the paper (<https://arxiv.org/pdf/1409.1556.pdf>).

What percentage of these parameters are due to the fully connected layers?

Compare this with the number of parameters in a network with many more layers like ResNet or GoogLeNet. These networks although deeper have a lot fewer parameters. Comment why?

Note: We are only interested in the rough estimate (ex: 1 million, 100k etc) and not the exact value(ex: 123,456).

6 Submission

6.1 Handin Zip Structure

- <andrew id>/
 - <andrew id>.pdf
 - code/
 - Backwards.m (30 points)
 - checkGradient.m (25 points)
 - Classify.m (5 points)
 - ComputeAccuracyAndLoss.m (10 points)
 - extractImageText.m (10 points)
 - findLetters.m (20 points)
 - finetune36.m (5 points)
 - Forward.m (10 points)
 - InitializeNetwork.m (5 points)
 - Train.m (10 points)
 - testFindLetters.m (5 points)
 - testExtractImageText.m (5 points)
 - train26.m (5 points)
 - UpdateParameters.m (5 points)
 - Any other code or helper functions you used
- ec/ *optional*
 - README
 - Single test script for each attempted extra credit
 - Any other extra credit code

6.2 Writeup Summary

Q1.1.1 Theory [5 points]] ReLU?

Q1.1.2 Theory [5 points]] Linear activation function.

Q2.1.1 Theory [5 points]] Zero or one initialization.

Q2.1.3 Writeup [5 points]] Describe initialization strategy.

Q2.4.1 Theory [5 points]] Batch vs. stochastic gradient descent.

Q3.1.2 Writeup [5 points]] Training plots and comments on `nist26` data

Q3.1.3 Writeup [5 points]] Network test set performance and weights visualization

Q3.1.4 Writeup [5 points]] Confusion matrix image visualization and comments

Q3.2.1 Writeup [5 points]] Training plot for fine tuning

Q3.2.2 Writeup [5 points]] Fine tuned weights visualization

Q3.2.3 Writeup [5 points]] Confusion matrix image for fine tuned network

Q4.1 Theory [5 points]] Failure cases for sample character detection algorithm

Q4.2 Writeup [5 points]] Character detection results on each of the images in `images/`

Q4.3 Writeup [5 points]] Extracted text from each of the images in `images/`

References

- [1] P. J. Grother. Nist special database 19 handprinted forms and characters database. <https://www.nist.gov/srd/nist-special-database-19>, 1995.
- [2] Xavier Glorot and Yoshua Bengio. Understanding the difficulty of training deep feed-forward neural networks.
- [3] John Duchi, Elad Hazan, and Yoram Singer. Adaptive subgradient methods for on-line learning and stochastic optimization. *Journal of Machine Learning Research*, 12(Jul):2121–2159, 2011.
- [4] Matthew D Zeiler. Adadelata: an adaptive learning rate method. *arXiv preprint arXiv:1212.5701*, 2012.
- [5] Diederik Kingma and Jimmy Ba. Adam: A method for stochastic optimization. *arXiv preprint arXiv:1412.6980*, 2014.